

AI Developer — Test Assignment

Goal

Build and deliver a small working project that demonstrates how you design and implement an AI Agent that solves a realistic business case end-to-end.

Deliverable

1. A public Git repository link (GitHub/GitLab).
2. A clear README that includes:
 - Problem statement (business context)
 - What the agent does (scope + non-scope)
 - Business cases covered (list)
 - Architecture overview (components + data flow)
 - Setup instructions (local run)
 - How to test/verify (unit/integration tests + manual сценарії)
 - Example inputs/outputs (screenshots/logs are fine)
 - Trade-offs, limitations, and next steps
 - Whether this solution is used in a real business today (see “Real-world usage” section below)

Core requirement

Your project must implement at least one business workflow using an AI Agent approach (planner/executor, tool-using agent, function-calling, multi-step reasoning with guardrails).

You may use OpenClaw or any alternative agent framework (e.g., LangGraph, OpenAI Agents SDK, Semantic Kernel, custom orchestrator). Choose what you prefer—just justify the choice.

Minimum functional scope (must have)

- Agent orchestration: multi-step workflow (not a single prompt)
- Tools/functions: at least 2 tools (e.g., search over local docs, CRUD store, email drafting, validation, classifier, web fetch mock etc)
- Reliable outputs: structured JSON result + human-readable summary
- Error handling and guardrails: retries, invalid input handling, refusal rules where appropriate
- Observability: logs/traces of tool calls and final decision summary
- Tests: at least basic unit tests + 1 integration/e2e test

Evaluation criteria

- Product thinking: is the business case clearly defined and solved?
- Agent design quality: planning, tool use, state management, prompt hygiene
- Engineering quality: code structure, typing, tests, readability, skills, prompt injecting sec

- Reliability: edge cases, deterministic checks, validation
- Clarity: README and instructions, easy to run and verify
- Practicality: could a small team extend this into production?

Real-world usage (required section in README)

Add a short section answering:

- “Is this solution currently used in any business?”
 - If yes: describe where and how (high level, no confidential info).
 - If no: state “Not in production today” and describe 2–3 realistic businesses/teams that could use it, plus what would be required to productionize (security, monitoring, data access, human-in-the-loop, compliance).

Submission

Send the Git repository link. Ensure the repository includes:

- Source code
- README with setup + verification steps
- Sample inputs (and expected outputs)
- Tests and instructions to run them